

Sender-Initiated

The sender-initiated load distribution algorithm is a technique used in distributed systems to balance the load across multiple nodes or servers. In this algorithm, the sender (client or source) is responsible for distributing the workload among available nodes.

Here's a high-level description of how the sender-initiated load distribution algorithm works:

1. **Sender sends a request:** The sender initiates a request to perform a task or process a workload. This sender can be a client application or a source node in the distributed system.
2. **Load information:** The sender gathers information about the current load or availability of nodes in the system. This information can include metrics such as CPU utilization, memory usage, network traffic, or any other relevant parameters that indicate the workload of each node.
3. **Load balancing decision:** Based on the load information, the sender makes a load balancing decision to determine which node should receive the request. The decision can be based on various algorithms and policies, such as round-robin, weighted round-robin, least-connection, or any custom logic that considers the current load and availability of nodes.
4. **Assigning the task:** Once the load balancing decision is made, the sender assigns the task or workload to the selected node. This can involve sending the request directly to the chosen node or using a routing mechanism to forward the request.
5. **Monitoring and adaptation:** After assigning the task, the sender continues to monitor the system's load and performance. If significant changes occur in the workload distribution or if a node becomes unavailable or overloaded, the sender may adapt its load balancing decision and reassign the tasks to different nodes to maintain balanced load distribution.
6. **Feedback and coordination:** The sender may receive feedback or coordination from the nodes regarding their current load, availability, or completion status of the assigned tasks. This feedback can help the sender in making informed load balancing decisions and adjusting the workload distribution as necessary.
7. **Repeat steps 2-6:** The sender-initiated load distribution algorithm continues to operate in a loop, continuously monitoring the load, making load balancing decisions, assigning tasks, and adapting to changes in the system's load and availability.

By using the sender-initiated load distribution algorithm, the sender takes an active role in load balancing and distributes the workload based on its assessment of the

system's load. This approach allows for flexibility and adaptability in load distribution, as the sender can make real-time decisions based on dynamic workload conditions.

It's worth mentioning that the effectiveness of the sender-initiated load distribution algorithm depends on accurate load information and timely decision-making by the sender. Additionally, the algorithm's performance can be influenced by network latency, scalability of the load balancing logic, and the overall system architecture.

While the sender-initiated load distribution algorithm offers several benefits, it also has some disadvantages. Here are a few:

1. Increased network traffic: With the sender taking responsibility for load distribution, there is an increased amount of communication between the sender and the nodes. This can result in higher network traffic and potential congestion, especially when the system is dealing with a large number of requests or tasks.
2. Sender as a single point of failure: In this algorithm, the sender plays a crucial role in load balancing decisions. If the sender fails or becomes overloaded, it can become a single point of failure, affecting the entire load distribution process. This centralized dependency on the sender can decrease the system's fault tolerance and resilience.
3. Delay in load balancing decisions: Since the sender needs to gather load information, make load balancing decisions, and assign tasks, there can be a delay in the load distribution process. This delay can impact the system's ability to quickly adapt to changes in workload or node availability.
4. Limited scalability: The sender-initiated load distribution algorithm may face scalability challenges as the system grows. As the number of nodes increases, the sender's workload of collecting load information and making load balancing decisions can become more complex and time-consuming. This can impact the algorithm's efficiency and scalability.
5. Lack of real-time load information: The load information collected by the sender may not always reflect real-time conditions. By the time the sender makes load balancing decisions and assigns tasks, the load distribution situation may have changed. This lag in obtaining and processing load information can lead to suboptimal load balancing decisions.
6. Difficulty in handling heterogeneous nodes: In distributed systems, nodes may have different capabilities and resources. The sender-initiated load distribution algorithm may struggle to effectively handle heterogeneous nodes with varying processing power, memory, or network bandwidth. Load balancing decisions that do not account for these differences may result in uneven workload distribution.

To mitigate these disadvantages, it's important to carefully design and implement the sender-initiated load distribution algorithm, considering the specific requirements and characteristics of the distributed system. Alternative load distribution approaches, such as receiver-initiated algorithms or hybrid approaches, can also be explored to address some of the limitations mentioned above.

Receiver-Initiated

Here is a high-level algorithm for the receiver-initiated load distribution:

1. Initialize the distributed system with a set of receivers or nodes.
2. Each receiver periodically monitors its own workload and resource availability. This can be done by collecting local metrics such as CPU utilization, memory usage, or the number of pending tasks.
3. Based on its workload and resource availability, each receiver independently decides whether it needs to request additional tasks from the sender or if it can handle more tasks.
4. If a receiver determines that it requires additional tasks, it sends a load request message to the sender. The message typically includes information about the receiver's workload and resource availability.
5. The sender receives load request messages from the receivers and maintains a queue of pending tasks.
6. Upon receiving a load request message, the sender evaluates the workload and resource availability of the requesting receiver. It compares this information with the workload distribution across the receivers and determines whether to assign tasks to the requesting receiver.
7. If the sender decides to assign tasks, it selects a suitable number of tasks from the pending task queue and sends them to the requesting receiver.
8. The requesting receiver receives the assigned tasks from the sender and starts processing them.
9. Steps 2-8 are repeated periodically or whenever there are significant changes in the workload or resource availability of the receivers.
10. If a receiver becomes overloaded or experiences resource constraints, it can notify the sender and request to offload or redistribute some of its tasks to other receivers.
11. If a receiver or the sender fails, the remaining receivers continue to operate and distribute the load among themselves. Failed nodes can be detected through heartbeats or other failure detection mechanisms.
12. The algorithm continues to dynamically balance the workload across the receivers, adapt to changes in workload and resource availability, and maintain an efficient load distribution within the distributed system.

It's important to note that the specific implementation details of the receiver-initiated load distribution algorithm can vary depending on the characteristics of the distributed system and the requirements of the application. The algorithm described above provides a high-level overview of the main steps involved in the receiver-initiated approach.

Disadvantages:

1. Increased communication overhead: The receiver-initiated algorithm requires constant communication between the receivers and the sender to request and receive tasks. This can introduce additional communication overhead, especially in large-scale distributed systems, potentially affecting network performance.
2. Load imbalance potential: In the receiver-initiated algorithm, the workload distribution depends on each receiver's individual decision-making. This can lead to load imbalances if certain receivers consistently request more tasks than others or if the workload distribution is not optimized across the system.
3. Complexity in load estimation: Nodes must estimate their own workload accurately to make effective task requests. Load estimation can be challenging, especially in dynamic environments where workload fluctuates rapidly. Inaccurate load estimation can lead to inefficient load balancing decisions.
4. Dependency on node cooperation: The effectiveness of the receiver-initiated algorithm relies on the cooperation and reliability of the individual nodes. If some nodes are unresponsive or unreliable, it can impact the load balancing performance and system efficiency.
5. Limited global knowledge: Unlike the sender-initiated algorithm, where the sender has a global view of the system's load, the receiver-initiated algorithm operates based on local information available to each receiver. This limited global knowledge can make it harder to achieve an optimal load balancing solution across the entire system.

Overall, the receiver-initiated load distribution algorithm offers benefits such as decentralized decision-making and dynamic load adaptation. However, it also introduces challenges related to communication overhead, load imbalance, load estimation, node cooperation, and limited global knowledge. The choice between sender-initiated and receiver-initiated algorithms depends on the specific requirements and characteristics of the distributed system.

Symmetric

- Senders search for receivers and vice-versa.

- Low loads: senders can find receivers easily. High loads: receivers can find senders easily.
- May have disadvantages of both: polling at high loads can make the system unstable. Receiver-initiated task transfers can be preemptive and so expensive.
- Simple algorithm: combine previous two approaches.
- Above-average algorithm:
 - Transfer Policy: Two adaptive thresholds instead of one. If a node's estimated average load is A , a higher threshold $\text{TooHigh} > A$ and a lower threshold $\text{TooLow} < A$ are used.
 - $\text{Load} < \text{TooLow} \rightarrow \text{receiver}$. $\text{Load} > \text{TooHigh} \rightarrow \text{sender}$
- Sender Component
 - Node with TooHigh load, broadcasts a TooHigh message, sets TooHigh timer, and listens for an Accept message.
 - A receiver that gets the (TooHigh) message sends an Acceptmessage, increases its load, and sets AwaitingTask timer.
 - If the AwaitingTask timer expires, load is decremented.
 - On receiving the Accept message: if the node is still a sender, it chooses the best task to transfer and transfers it to the node.
 - When sender is waiting for Accept, it may receive a TooLow message (receiver initiated). Sender sends TooHigh to that receiver. Do step 2 & 3.
 - On expiration of TooHigh timer, if no Accept message is received, system is highly loaded. Sender broadcasts a ChangeAverage message.
- Receiver Component
 - Node with TooLowload, broadcasts a TooLow message, sets a TooLow timer, and listens for TooHigh message.
 - If TooHighmessage is received, do step 2 & 3 in Sender Component.
 - If TooLow timer expires before receiving any TooHigh message, receiver broadcasts a ChangeAveragemessage to decrease the load estimate at other nodes.

Selecting an Algorithm

- If a system never gets highly loaded, sender-initiated algorithms work better.
- Stable, receiver-initiated algorithms better for high loads.
- Widely fluctuating loads: stable, symmetric algorithms.
- Widely fluctuating loads + high migration cost for preemptive transfers: stable, sender-initiated algorithms.
- Heterogeneous work arrival: stable, adaptive algorithms.