VIT®

Vellore Institute of Technology

(Deemed to be University under section 3 of UGC Act, 1956)

ITA-6008 Big Data Analytics

Assignment - 1

By :

Rajat Singh 22MCA0139

Submitted to :

Prof. POUNAMBAL M

1. Procedure to install the Hadoop in your system.

Prerequisite to Hadoop Installation

1. You have installed Ubuntu 22 Desktop version in your Virtual Machine
2. You have installed Java (jdk 8) in your Ubuntu system.
3. Check your hostname is Ubuntu

$ hostname --should output Ubuntu

Linux Configuration Before Hadoop Installation

We will setup single node Hadoop cluster using a dedicated Hadoop user.

1. Login as Root

2. Adding a dedicated user called hduser

2. Create a Group called Hadoop

```
sudo addgroup Hadoop
```

4. Create an User hduser

```
sudo adduser hduser
```

5. Add houser to hadoop group

```
sudo adduser houser Hadoop
```

6. Add the 'hduser' to sudoers list so that hduser can do admin tasks.
sudo visudo
houser ALL=(ALL) ALL

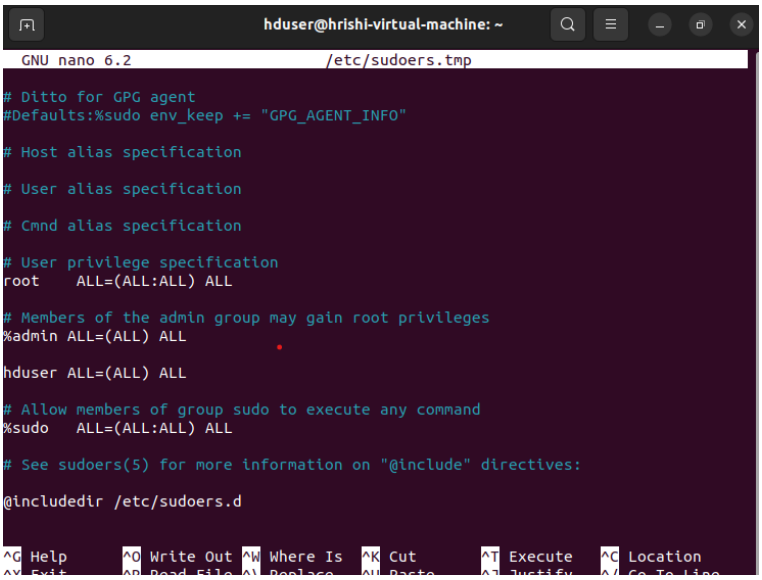

7. Logout Your System and login again as hduser.

8. Configuring SSH

```
sudo apt-get install openssh-server
```

9. Generate SSH for communication

```
hduser@ubuntu:~$ ssh-keygen
```

10.Copy Public Key to Authorized key file & edit the permission

```
hduser@ubuntu:~Scat~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

```
ho&ser@ubuntu:~$chmod 700 ~/.ssh/auth
```

11.Start SSH

If ssh is not running, then run it by giving the below command

```
hduser@ubuntu:~$ sudo /etc/init.d/ssh restart
```

Enter your Password(hadoop)

12. Test Your SSH Connectivity

```
hduser@ubuntu:~$ ssh localhost
```

13. Test Your SSH Connectivity

```
hduser@ubuntu:-S ssh localhost
```

Download Hadoop

1. Download hadoop-3.3.1.tar.gz and save it to hduser/Desktop.
https://downloads.apache.org/hadoop/

2. move the above downloaded file to /us/local/
Open Terminal(Ctr|+Alt+T)

```
$ sudo mv ~/Desktop/hadoop-3.3.1.tar.gz /us/local/
cd /usr/local
sudo tar -xvf hadoop-3.3.1.tar.gz
sudo rm hadoop-3.3.1.tar.gz
sudo ln -s hadoop-3.3.1 hadoop
sudo chown -R hduser:hadoop hadoop-3.3.1
sudo chmod 777 hadoop-3.3.1
```

3. Edit hadoop-env.sh and configure Java.

```
$ sudo vim /ust/local/hadoop/etc/hadoop/hadoop-env.sh
```

```
export HADOOP_OPTS=Djava.net.preferiPv4Stack=true
export HADOOP_HOME_WARN_SUPPRESS-"TRUE"
export JAVA_HOME=/us/local/java/jdk
```

3. Update $HOME/.bashrc

```
# Set Hadoop-related environment variables
export HADOOP_HOME=/usr/local/hadoop
export HADOOP_MAPRED_HOME=${HADOOP_HOME}
export HADOOP_COMMON_HOME=${HADOOP_HOME}
export HADOOP_HDFS_HOME=${HADOOP_HOME}
export HADOOP_YARN_HOME=${HADOOP_HOME}
export HADOOP_CONF_DIR=${HADOOP_HOME}/etc/hadoop

# Native Path
export
HADOOP_COMMON_LIB_NATIVE_DIR=${HADOOP_PREFIX}/lib/native
export HADOOP_OPTS="-Djava.library.path=$HADOOP_PREFIX/lib"

# Set JAVA_HOME (we will also configure JAVA_HOME directly for Hadoop later
on)
export JAVA_HOME=/usr/local/java/jdk
# Some convenient aliases and functions for running Hadoop-related commands
unaliasfs&> /dev/null
aliasfs="hadoop fs"
unaliashls&> /dev/null
aliashls="fs -ls"
```

export
PATH=$PATH:$HADOOP_HOME/bin:$PATH:$JAVA_HOME/bin:$HADOOP_HOME/sbin

5. Update yarn-site.xml

$sudo vim /us/local/hadoop/etc/hadoop/yarn-site.xmI

```
<property>
        <name>yarn.nodemanager.aux-services</name>
        <value>mapreduce_shuffle</value>
    </property>
    <property>
  <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
  <value>org.apache.hadoop.mapred.ShuffleHandler</value>
    </property>
  <property>
    <name>yarn.nodemanager.vmem-check-enabled</name>
    <value>false</value>
    <description>Whether virtual memory limits will be enforced for containers</description>
  </property>

  <property>
    <name>yarn.nodemanager.vmem-pmem-ratio</name>
    <value>4</value>
    <description>Ratio between virtual memory to physical memory when setting memory limits for containers</description>
  </property>
```

6. Update core-site.xml file

$ sudo vim /ust/local/hadoop/etc/hadoop/core-site.xmI

```
<property>
        <name>hadoop.tmp.dir</name>
        <value>/app/hadoop/tmp</value>
        <description>A base for other temporary directories.</description>

    </property>

 <property>
     <name>fs.default.name</name>
    <value>hdfs://localhost:9000</value>
    <description>default host and port</description>
    </property>
```

```
<property>
  <name>hadoop.proxyuser.hduser.hosts</name>
  <value>*</value>
 </property>

<property>
  <name>hadoop.proxyuser.hduser.groups</name>
  <value>*</value>
</property>
```

7. Create the above temp folder and give appropriate permission

```
sudo mkdir -p /app/hadoop/tmp
sudo chown hduser:hadoop -R /app/hadog
sudo chmod 750 /app/hadoop/tmp
```

8. Edit mapred-site.xml

```
sudo vim /us/local/hadoop/etc/hadoop/mapred-site.xml
```

```
<property>
  <name>mapreduce.framework.name</name>
      <value>yarn</value>
      </property>
 <property>
      <name>mapreduce.jobhistory.address</name>
          <value>localhost:10020</value>
          <description>Host and port for Job History Server (default

          0.0.0.0:10020)</description>
</property>

  <property>
    <name>yarn.app.mapreduce.am.env</name>
    <value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>
   </property>
   <property>
    <name>mapreduce.map.env</name>
    <value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>
   </property>
   <property>
    <name>mapreduce.reduce.env</name>
    <value>HADOOP_MAPRED_HOME=${HADOOP_HOME}</value>
   </property>
```

9. Create a temporary directory which will be used as base location for

```
sudo mkdir -p /us/local/hadoop_tmp/hdfs/namenode
sudo,mkdir -p /ust/local/hadoop_tmp/hdfs/datanode
sudo chown hduser:hadoop -R /us/local/hadoop tmp/
```

10. Update hdfs-site.xmI file

```
$ sudo vim /us/local/hadoop/etc/hadoop/hdfs-site.xml
```

```
<property>
  <name>dfs.replication</name>
  <value>1</value>
 </property>
 <property>
  <name>dfs.namenode.name.dir</name>
  <value>file:/usr/local/hadoop_tmp/hdfs/namenode</value>
 </property>
 <property>
  <name>dfs.datanode.data.dir</name>
  <value>file:/usr/local/hadoop_tmp/hdfs/datanode</value>
 </property>
```

11.Format your namenode

```
$ hadoop namenode -format
```

12.Starting your single-node cluster

```
$ start-all.sh
```

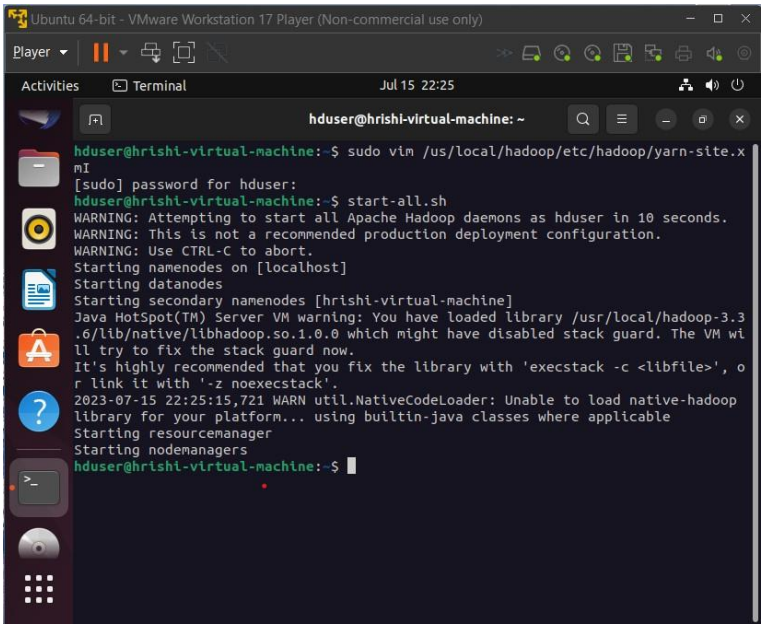

13. Start your history-server.

```
$ mr-jobhistory-daemon.sh start historyserver
$ mr-jobhistory-daemon.sh stop historyserver
```

14. Check if all the necessary hadoop daemon is running or not

```
$ jps
```

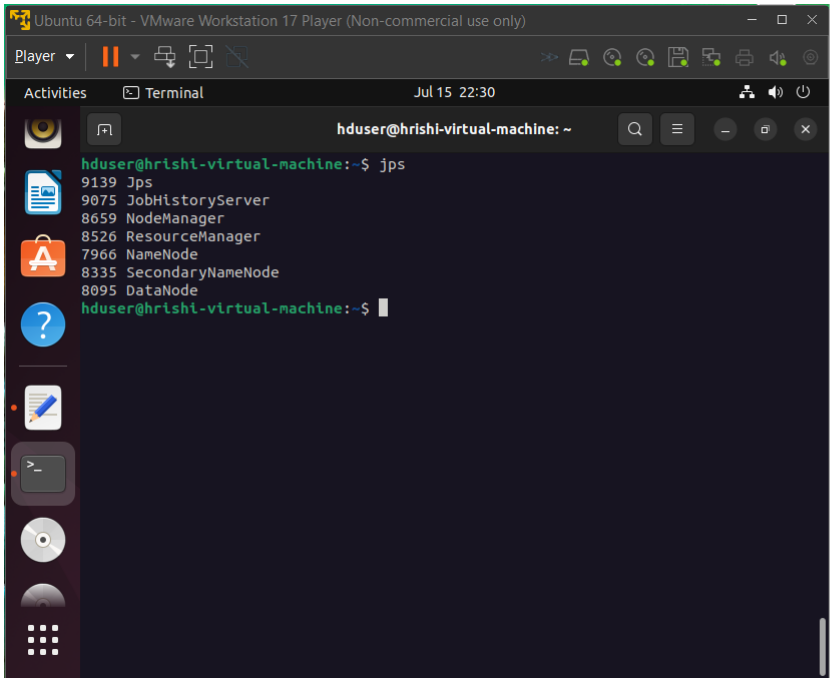

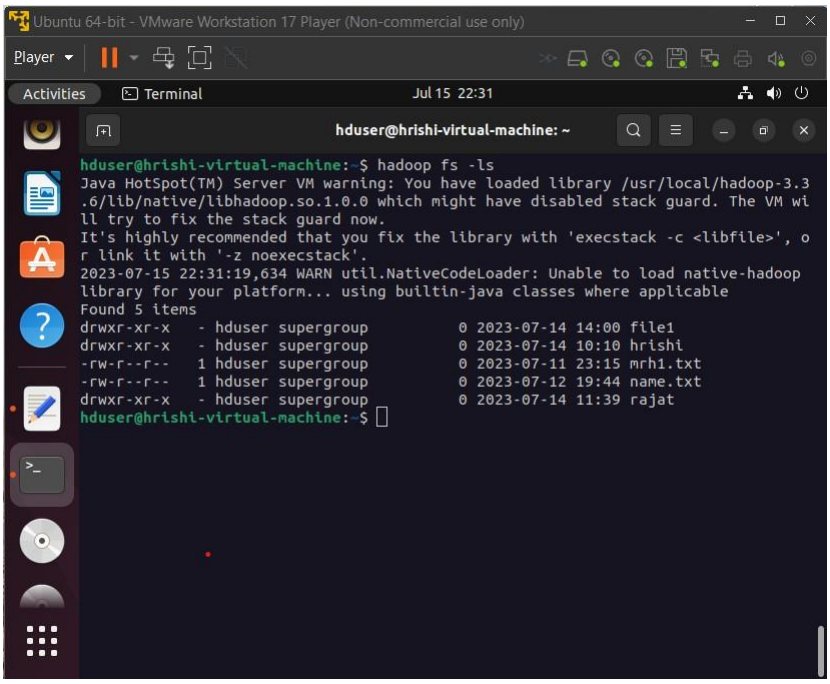

2. Creation of sample table in HABASE with 7 attributes with 2 sub columns and displaying of valid data.

Table Creation :

create 'student1', 'name', 'regno', 'year'', 'age', 'address', 'mobile', 'marks'

```
hbase(main):006:0> create 'student1','name','regno','year','age','address','mobile','marks'
0 row(s) in 2.3200 seconds

=> Hbase::Table - student1
hbase(main):007:0>
```

put ‘student1’, 'row1', 'Name', 'John'

put ' student1', 'row1', 'RegNo', '2021001'

put 'student1', 'row1', 'Year', '2023'

put ‘student1’, 'row1', 'Age', '25'

put ‘student1’, 'row1', 'Mobile', '9876543210'

put ‘student1’, 'row1', 'Address', '123 Main Street'

put ‘student1’, 'row1', 'Marks:Marks1', '80'

put ‘student1’, 'row1', 'Marks:Marks2', '75'

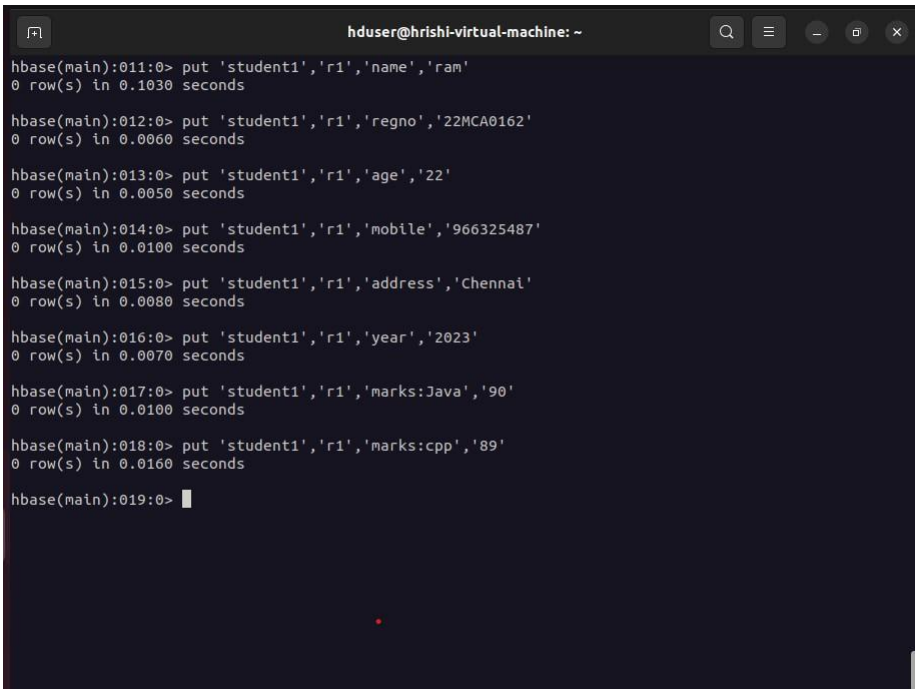
hduser@hrishi-virtual-machine: ~
hbase(main):011:0> put 'student1','r1','name','ram'
0 row(s) in 0.1030 seconds
hbase(main):012:0> put 'student1','r1','regno','22MCA0162'
0 row(s) in 0.0060 seconds
hbase(main):013:0> put 'student1','r1','age','22'
0 row(s) in 0.0050 seconds
hbase(main):014:0> put 'student1','r1','mobile','966325487'
0 row(s) in 0.0100 seconds
hbase(main):015:0> put 'student1','r1','address','Chennai'
0 row(s) in 0.0080 seconds
hbase(main):016:0> put 'student1','r1','year','2023'
0 row(s) in 0.0070 seconds
hbase(main):017:0> put 'student1','r1','marks:Java','90'
0 row(s) in 0.0100 seconds
hbase(main):018:0> put 'student1','r1','marks:cpp','89'
0 row(s) in 0.0160 seconds
hbase(main):019:0>

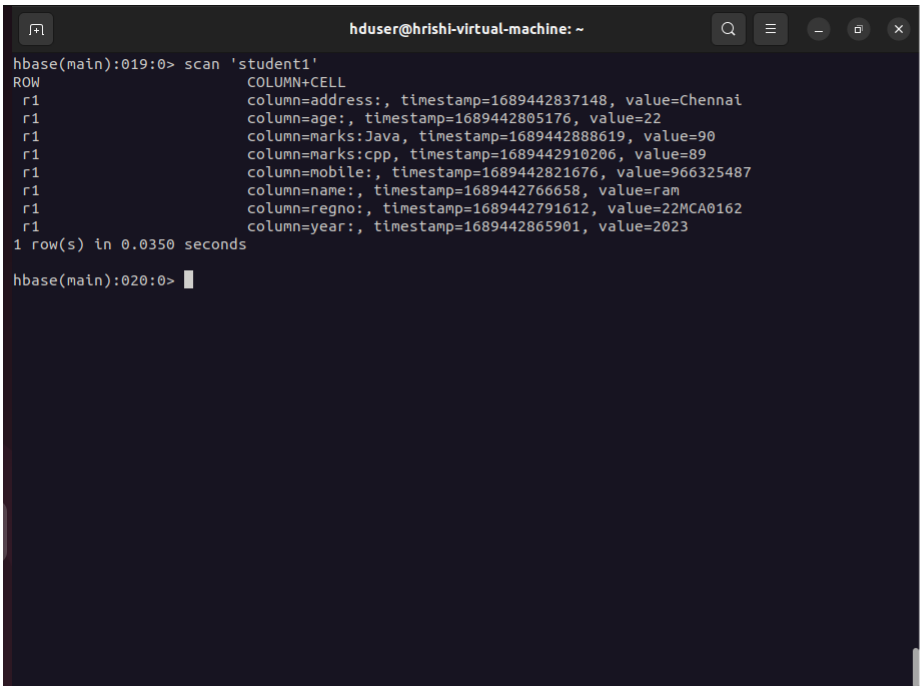
hduser@hrishi-virtual-machine: ~
hbase(main):019:0> scan 'student1'
ROW COLUMN+CELL
r1 column=address:, timestamp=1689442837148, value=Chennai
r1 column=age:, timestamp=1689442805176, value=22
r1 column=marks:Java, timestamp=1689442888619, value=90
r1 column=marks:cpp, timestamp=1689442910206, value=89
r1 column=mobile:, timestamp=1689442821676, value=966325487
r1 column=name:, timestamp=1689442766658, value=ram
r1 column=regno:, timestamp=1689442791612, value=22MCA0162
r1 column=year:, timestamp=1689442865901, value=2023
1 row(s) in 0.0350 seconds
hbase(main):020:0>